A Project Report

On

IR system for Policy Documents

BY

Under the supervision of

Dr. Aruna Malapati

SUBMITTED IN PARTIAL FULLFILLMENT OF THE REQUIREMENTS OF

CS F469: Information Retrieval

BIRLA INSTITUTE OF TECHNOLOGY AND SCIENCE PILANI (RAJASTHAN)

HYDERABAD CAMPUS

(March 2023)

Group Details

Name	ID no.
Aaditya Mahesh Rathi	2020A7PS2191H
Sankalp Kulkarni	2020A7PS1097H
Akshat Oke	2020A7PS0284H

ACKNOWLEDGEMENTS

Table of Contents

Contents

1. Pipeline

1.1 Data Pre-processing

Initially, we utilized PyPDF2 tool to transform PDFs into text files. However, it resulted in merged words due to the removal of whitespaces. We tried to search for a way for decompounding the words using dictionary de compounders, but unfortunately it was only available for German language. After consulting with Dr. Aruna Malapati Ma'am, she said we can also use web-based converters for the same which solved this issue. We assumed that a single blank line separated paragraphs, but it resulted in small paragraphs. Hence, we set a minimum length of 20 words for each paragraph. If a paragraph was shorter than 20 words, it was appended to the next paragraph.

1.2 Document Expansion

Document Expansion has turned out to be an effective strategy in sparse retrieval. The effectiveness of Document Expansion by Query Prediction was first proven in [1]. Since then, it has been used in multiple state-of-the-art approaches for sparse retrieval tasks. Many approaches that came out after this paper and achieved better results were based on the same idea. [2] [3]

[1] used the Doc2Query model to predict queries. It was a sequence-to-sequence model. But we used DocT5query, which is a transformer-based model, to generate questions for the paragraphs. Since paragraph size was comparatively small for the dataset given to us, we only predicted 1 query for each document (in this scenario paragraph). We used the implementation of DocT5query provided in the standard transformer library.

1.3 Tokenization and Stemming

Initially, we considered spacy and nltk for Tokenization and Stemming. However, we discovered that spacy was slow and performed tasks such as NER annotation and PoS tagging, which were irrelevant to our project. Additionally, spacy used 326 stop words, whereas nltk used only 179 stop words. As a result, we decided to avoid Tokenization and Stemming using spacy and instead used the standard nltk tokenizer. nltk had two stemmers for English language- Lancaster and Porter. Lancaster stemmer, however, tends to

over-stem words aggressively in many cases. Therefore, we opted for Porter stemmer for stemming.

1.4 Index Construction

Our IR system supports free text queries and phrase queries. Based on our analysis of dataset, we decided that wild-card queries weren't relevant. Also, computational cost for it wasn't worth the benefit it would provide. We built two types of indexes. Inverted Index for normal free text queries and Bigram Index for phrase queries. Based on our dataset, we anticipated that most phrase queries would consist of two or three words, so we decided to use the bigram index. Since size of our dataset was so small, it could easily fit into the main memory, so we avoided using data structures like linked lists or B-trees, which would have slowed down our system with unnecessary memory references. Using such data structures is useful when the data doesn't fit entirely in the main memory and we need to fetch multiple blocks of the disk. We used Python lists to store the posting lists. Python arrays are just wrappers around C arrays and are more memory efficient, but Python lists are highly optimized and provide greater flexibility. They are recommended to be used in almost all cases other than where compatibility with C code is necessary. We used a dictionary to store terms and their corresponding posting lists.

1.5 Ranking

We used BM25 algorithm for ranking the retrieved documents. TF-IDF penalizes document frequency and favors term frequency. It doesn't consider the issue of term saturation. If a document contains a word 50 times and another one contains it 100 times, is it twice as relevant as the others? In most cases, both documents would be equally important. If document is shorter and contains the term then it is more likely to be relevant than a longer document that contains the term, as it is likely to mention multiple things. Below is the formula for the BM25 ranking function.

$$\sum_i^n IDF(q_i) \frac{f(q_i, D) * (k1 + 1)}{f(q_i, D) + k1 * (1 - b + b * \frac{fieldLen}{avgFieldLen})}$$

b and k are hyperparameters here. b controls effect of length of document in its score. If b is larger, documents with longer length than average document length are penalized more (scored less). Literature indicates that value of b = 0.75 works well in most IR tasks. k1 controls term saturation. Higher k1 indicates that saturation value for term frequency is higher. Increase in score when term frequency grows beyond certain value will be minimal. Literature shows that k1 = 1.2 works well in most IR tasks. We used b = 1 and k1=0.5. Since the document length in our case was less, the term frequency was also less, so we decided to use these values.

2. References

References

[1] W. Y. J. L. K. C. Rodrigo Nogueira, "Document Expansion by Query Prediction," 2019.

[2] C. L. B. P. S. C. Thibault Formal, "SPLADE v2: Sparse Lexical and Expansion Model for Information Retrieval," 2021.

[3] J. C. Zhuyun Dai, "Context-Aware Sentence/Passage Term Importance Estimation For First Stage Retrieval," 2019.